

dl_3

April 11, 2024

```
[1]: import tensorflow as tf
      from tensorflow import keras
      import numpy as np
      import matplotlib.pyplot as plt
```

C:\Users\Nimisha jadhav\anaconda3\lib\site-packages\scipy__init__.py:155:
UserWarning: A NumPy version >=1.18.5 and <1.25.0 is required for this version
of SciPy (detected version 1.26.4
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")

```
[3]: fashion_mnist= keras.datasets.fashion_mnist
      (train_images, train_labels), (test_images, test_labels)= fashion_mnist.
      ↪load_data()
      #normalize the images
      train_images= train_images / 255.0
      test_images= test_images / 255.0
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>
29515/29515 0s 5us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>
26421880/26421880 40s
2us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>
5148/5148 0s 23us/step
Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>
4422102/4422102 12s
3us/step

```
[5]: #define the model
      model= keras.Sequential([keras.layers.Conv2D(64,(3,3),
      ↪activation='relu',input_shape=(28,28,1)),
      keras.layers.MaxPooling2D(2,2),
      keras.layers.Flatten(),
      keras.layers.Dense(128,activation='relu'),
```

```
keras.layers.Dense(10,activation='softmax'))]
```

C:\Users\Nimisha jadhav\anaconda3\lib\site-packages\keras\src\layers\convolutional\base_conv.py:99: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__()
```

```
[9]: # Compile the model with the correct loss function
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
             metrics=['accuracy'])
```

```
[10]: #train the model
model.fit(train_images, train_labels, epochs=10)
#evaluate the model
test_loss, test_acc=model.evaluate(test_images, test_labels)
print('Test accuracy', test_acc)
```

```
Epoch 1/10
1875/1875          34s 18ms/step -
accuracy: 0.8241 - loss: 0.5051
Epoch 2/10
1875/1875          36s 19ms/step -
accuracy: 0.9054 - loss: 0.2589
Epoch 3/10
1875/1875          37s 20ms/step -
accuracy: 0.9237 - loss: 0.2059
Epoch 4/10
1875/1875          36s 19ms/step -
accuracy: 0.9338 - loss: 0.1734
Epoch 5/10
1875/1875          36s 19ms/step -
accuracy: 0.9467 - loss: 0.1441
Epoch 6/10
1875/1875          36s 19ms/step -
accuracy: 0.9549 - loss: 0.1226
Epoch 7/10
1875/1875          36s 19ms/step -
accuracy: 0.9637 - loss: 0.1011
Epoch 8/10
1875/1875          36s 19ms/step -
accuracy: 0.9707 - loss: 0.0800
Epoch 9/10
1875/1875          36s 19ms/step -
accuracy: 0.9764 - loss: 0.0662
Epoch 10/10
1875/1875          36s 19ms/step -
```

```

accuracy: 0.9804 - loss: 0.0547
313/313          2s 5ms/step -
accuracy: 0.9179 - loss: 0.3547
Test accuracy 0.9182999730110168

```

```

[12]: #make predictions
predictions= model.predict(test_images)
predicted_labels= np.argmax(predictions, axis=1)

```

```

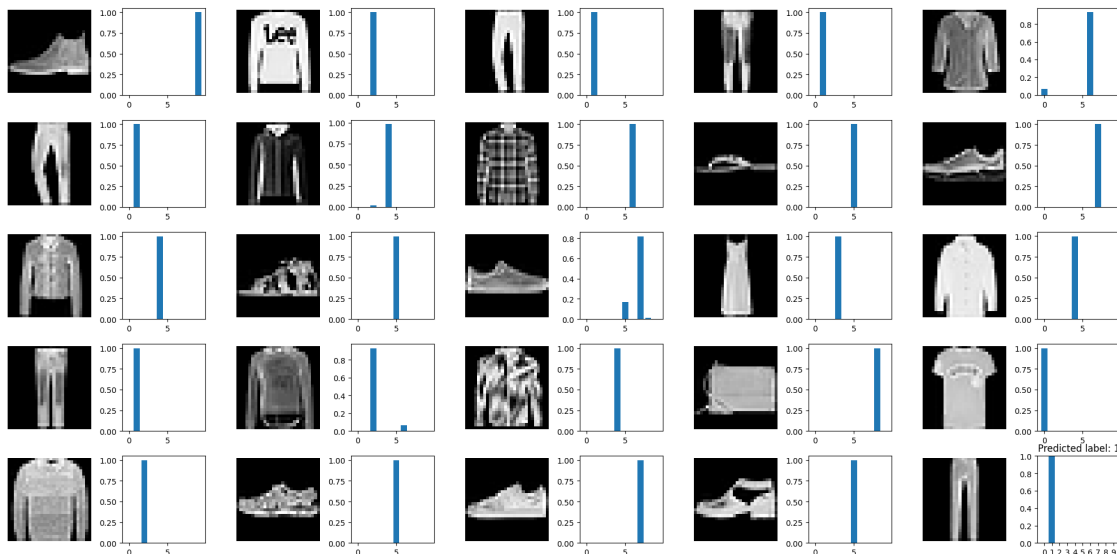
313/313          2s 5ms/step

```

```

[18]: num_rows=5
num_cols=5
num_images= num_rows * num_cols
plt.figure(figsize=(2 * 2 *num_cols, 2 * num_rows))
for i in range(num_images):
    plt.subplot(num_rows, 2* num_cols, 2 * i + 1)
    plt.imshow(test_images[i],cmap='gray')
    plt.axis('off')
    plt.subplot(num_rows, 2 * num_cols, 2* i + 2)
    plt.bar(range(10),
predictions[i])
plt.xticks(range(10))
plt.ylim([0, 1])
plt.tight_layout()
plt.title(f"Predicted label: {predicted_labels[i]}")
plt.show()

```



```

[ ]:

```